



DeepLearning.AI

Introduction to NLP with PyTorch

Working with text using PyTorch

In Module 3 you'll dive into:



Optimization



Images




Text




Efficiency


Images are represented as numerical values



Text has no intrinsic numeric meaning



Text has no intrinsic numeric meaning



Text has no intrinsic numeric meaning

Bob made a book collector happy the other day.



Last week, Bob made a book collector's day.

The other day, Bob delighted a book collector.

Bob recently brought delight to a book collector.

Text is sequential and contextual

Text is sequential and contextual

"A **bat** flew out of the cave"



Text is sequential and contextual


"A **bat** flew out of the cave"



"He swung the baseball **bat**"



NLP workflows require:



Specific steps

For text processing

NLP workflows require:



Specific steps

For text processing



Specialized models

To capture meaning
and context

NLP workflows require:



Specific steps

For text processing



Specialized models

To capture meaning
and context



Large datasets

To handle complexity

What makes text data challenging?

What makes text data challenging?



Sequence
dependencies
and context

What makes text data challenging?



Sequence
dependencies
and context



**Variable length
and structure**

What makes text data challenging?



Sequence
dependencies
and context



Variable length
and structure



Vocabulary and
representation
challenges

What makes text data challenging?



Sequence
dependencies
and context



Variable length
and structure



Vocabulary and
representation
challenges



Ambiguity and
polysemy

What makes text data challenging?



Sequence
dependencies
and context



Variable length
and structure



Vocabulary and
representation
challenges



Ambiguity and
polysemy

NLP applications: Classification

NLP applications: Classification



Sentiment analysis

NLP applications: Classification



Sentiment analysis



Spam detection

NLP applications: Classification



Sentiment analysis



Spam detection



Intent classification

NLP applications: Classification



Sentiment analysis



Spam detection



Intent classification



Topic categorization

NLP applications: Labeling

Named entity recognition (NER)

- Search engines
- Virtual assistants
- Medical record analyzers

NLP applications: Labeling

Named entity recognition (NER)

- Search engines
- Virtual assistants
- Medical record analyzers

Alice works at DeepLearning.AI in Mountain View.

The diagram illustrates the results of a Named Entity Recognition (NER) process on the sentence "Alice works at DeepLearning.AI in Mountain View.". The words "Alice", "DeepLearning.AI", and "Mountain View." are highlighted with colored boxes and have red arrows pointing down to their corresponding entity types: "person", "organization", and "location".

person organization location

NLP applications: Labeling

Named entity recognition (NER)

- Search engines
- Virtual assistants
- Medical record analyzers

Part-of-speech tagging

- Grammar checkers
- Speech recognition

NLP applications: Generative tasks

NLP applications: Generative tasks



Summarization

NLP applications: Generative tasks



Summarization



Dialogue generation

NLP applications: Generative tasks



Summarization



Dialogue generation



Machine translation

NLP applications: Generative tasks



Summarization



Dialogue generation



Machine translation



Creative writing

NLP applications: Generative tasks



Summarization



Dialogue generation



Machine translation



Creative writing

NLP applications: Generative tasks



Summarization



Dialogue generation



Machine translation



Creative writing


 PyTorch

The PyTorch logo consists of a stylized orange flame or swirl shape followed by the word "PyTorch" in a large, white, sans-serif font.

PyTorch is a popular framework for NLP



Preprocessing
text



Converting text to
tensors



Building models
for NLP tasks

PyTorch utilities for NLP

Before



TorchText: datasets,
tokenizers, iterators


PyTorch utilities for NLP

Before



TorchText: datasets,
tokenizers, iterators

Now



Tokenization
libraries and
pretrained models



DeepLearning.AI


Tokenization

Working with text using PyTorch

Tokenization: Dividing text into tokens

Tokenization: Dividing text into tokens

Words



Tokenization: Dividing text into tokens

Words

AI is the new electricity.

Subwords

The diagram illustrates the process of tokenizing words into subwords. On the left, a grey box labeled "Words" contains the sentence "AI is the new electricity.". To its right, another grey box labeled "Subwords" shows the same sentence broken down into subwords, each highlighted by a colored box and connected by arrows to the corresponding part of the original word. "AI" is shown as a single subword. "is" is shown as two subwords: "is". "the" is shown as one subword. "new" is shown as one subword. "electricity" is shown as three subwords: "electric-", "-", and "-ity".

AI
is
the
new
electricity


AI
is
the
new
electric-
-ity

Tokenization: Dividing text into tokens

Words

Subwords

Characters



Character tokenization

SILENT



S I L E N T

Character tokenization

SILENT



S I L E N T

=

LISTEN



L I S T E N

Numeric encodings: ASCII

Decimal	Character
64	@
65	A
66	B
67	C
68	D
...	...

Word tokenization

Word tokenization


“I love my dog”



I love my dog

Word tokenization

“I love my dog”



Special tokens

<UNK> → Unknown word

<PAD> → Pad sequences to match length

<START> → Beginning of sentence

<END> → End of sentence

Special tokens

<UNK> → Unknown word

<PAD> → Pad sequences to match length

<START> → Beginning of sentence

<END> → End of sentence

<CLS> → Classification inputs

<SEP> → Separate segments

Challenges with word tokenization



Vocabulary size

Challenges with word tokenization



Vocabulary size



Normalization

"Dog" vs. "dog"
"Hi!" vs. "Hi"
"don't" vs. "do not"

Building a tokenizer in PyTorch

```
sentences = [  
    'I love my dog',  
    'I love my cat'  
]  
  
# Tokenization function  
def tokenize(text):  
    # Lowercase the text and split by whitespace  
    return text.lower().split()
```

Building a tokenizer in PyTorch

```
sentences = [  
    'I love my dog',  
    'I love my cat'  
]  
  
# Tokenization function  
def tokenize(text):  
    # Lowercase the text and split by whitespace  
    return text.lower().split()
```

Building a tokenizer in PyTorch

```
sentences = [  
    'I love my dog',  
    'I love my cat'  
]  
  
# Tokenization function  
def tokenize(text):  
    # Lowercase the text and split by whitespace  
    return text.lower().split()
```

```
# Build the vocabulary
def build_vocab(sentences):
    vocab = {}
    # Iterate through each sentence.
    for sentence in sentences:
        # Tokenize the current sentence
        tokens = tokenize(sentence)
        # Iterate through each token in the sentence
        for token in tokens:
            # If the token is not already in the vocabulary
            if token not in vocab:
                # Add token to the vocabulary and assign it a unique integer ID
                # IDs start from 1; 0 can be reserved for padding.
                vocab[token] = len(vocab) + 1
    return vocab
```

```
# Build the vocabulary
def build_vocab(sentences):
    vocab = {}
    # Iterate through each sentence.
    for sentence in sentences:
        # Tokenize the current sentence
        tokens = tokenize(sentence)
        # Iterate through each token in the sentence
        for token in tokens:
            # If the token is not already in the vocabulary
            if token not in vocab:
                # Add token to the vocabulary and assign it a unique integer ID
                # IDs start from 1; 0 can be reserved for padding.
                vocab[token] = len(vocab) + 1
    return vocab
```

```
# Build the vocabulary
def build_vocab(sentences):
    vocab = {}
    # Iterate through each sentence.
    for sentence in sentences:
        # Tokenize the current sentence
        tokens = tokenize(sentence)
        # Iterate through each token in the sentence
        for token in tokens:
            # If the token is not already in the vocabulary
            if token not in vocab:
                # Add token to the vocabulary and assign it a unique integer ID
                # IDs start from 1; 0 can be reserved for padding.
                vocab[token] = len(vocab) + 1
    return vocab
```

Building a tokenizer in PyTorch

```
# Create the vocabulary index
vocab = build_vocab(sentences)

print("Vocabulary Index:", vocab)
```

Output

```
Vocabulary Index: {'i': 1, 'love': 2, 'my': 3, 'dog': 4, 'cat': 5}
```

Variable-length problem

Different sentences have different lengths




VS.

★★★★★ 5/28/2019

I have only written a couple of reviews in my life but after the amount of stress caused by my experience here I needed to say something. I came here for a cut and color. Originally a brunette, I had shared a detailed description of my hair history weeks prior to the apt, and images of the soft ash balayage I wanted as well as my current hair state. My hairdresser washed and cut my hair first, then dried it to apply color. From there I was explained that my hair was damaged and not able to create the look we had initially discussed. I understood, and happily agreed to something more subtle than we initially discussed (soft ashy natural brown balayage highlights was what we originally talked about). After the wash I was brought back to the station only to see a full transformation of my "damaged hair" from blonde to orange lemon blonde (that was not even toned to an ash, it was left orange). My hair was so bleached and broken, with patchy sections, and a far stretch from anything we discussed or agreed on. Im at a loss of understand how my hair was original damaged and unable to become what we agreed on but ok enough to be damaged far beyond anything we discussed. And a complete loss of what information I could have possibly given that would lend the dresser to think I wanted or requested to become a blonde. I paid and left (and even tipped). Only to scramble the next day trying to find someone to fix it before I had to meet my future mother in law for lunch. My hair is still falling out, and it's been a borderline tears and stress mixture since. I will not be coming back.

Variable-length problem

Padding



Adding <PAD> tokens
to match the longest
token in a batch

Variable-length problem

Padding



Adding <PAD> tokens
to match the longest
token in a batch

Truncation



Cutting longer
sequences down
to a fixed length

Padding

sequence length = 4

“I love my dog”

“I love”

Padding

sequence length = 4

“I love my dog”



[1, 2, 3, 4]

“I love”



[1, 2, <PAD>, <PAD>]

Truncation

sequence length = 2

"I love my dog"



"I love"



[1, 2]

Padding has a cost

“I love my dog, he is so playful and smart”



[1, 2, 3, 4, 5, 6, 7, 8, 9, 10]

“I love”



[1, 2, <PAD>, <PAD>, <PAD>, <PAD>, <PAD>, <PAD>, <PAD>, <PAD>, <PAD>]

Minimize the cost of padding

Bucketing



Group sequences of
similar length
together

Minimize the cost of padding

Bucketing



Group sequences of
similar length
together

Smart truncation



Preserve the
important parts of
a long text

Attention masks

Binary tensors that flag actual words to tell them apart from padding

Attention masks

Binary tensors that flag actual words to tell them apart from padding

Sentence:

“I love”



Tokens:

[1, 2, <PAD>, <PAD>]



Mask:

[1, 1, 0, 0]



DeepLearning.AI

Using a Pretrained Tokenizer

Working with text using PyTorch

NEW Get started with Inference in seconds! 



The AI community building the future.

The platform where the machine learning community collaborates on models, datasets, and applications.

[Explore AI Apps](#)

or [Browse 1M+ models](#)

Tasks Libraries Datasets Languages Licenses Other

Filter Tasks by name

Multimodal

 Text-to-Image  Image-to-Text
 Text-to-Video  Visual Question Answering
 Document Question Answering  Graph Machine Learning

Computer Vision

 Depth Estimation  Image Classification
 Object Detection  Image Segmentation
 Image-to-Image  Unconditional Image Generation
 Video Classification  Zero-Shot Image Classification

Natural Language Processing

 Text Classification  Token Classification
 Table Question Answering  Question Answering
 Zero-Shot Classification  Translation
 Summarization  Conversational
 Text Generation  Text2Text Generation
 Sentence Similarity

Audio

 Text-to-Speech  Automatic Speech Recognition
 Audio-to-Audio  Audio Classification
 Voice Activity Detection

Tabular

 Tabular Classification  Tabular Regression

Models 469,541

 Filter by name

meta-llama/Llama-2-70b

 Text Generation • Updated 4 days ago • ↓ 25.2k • ❤ 64

stabilityai/stable-diffusion-xl-base-0.9

Updated 6 days ago • ↓ 2.01k • ❤ 393

openchat/openchat

 Text Generation • Updated 2 days ago • ↓ 1.3k • ❤ 136

llyyasyvlie/ControlNet-v1-1

Updated Apr 26 • ❤ 1.87k

cerspense/zeroscope_v2_XL

Updated 3 days ago • ↓ 2.66k • ❤ 334

meta-llama/Llama-2-13b

 Text Generation • Updated 4 days ago • ↓ 328 • ❤ 64

tiiuae/falcon-40b-instruct

 Text Generation • Updated 27 days ago • ↓ 288k • ❤ 899

WizardLM/WizardCoder-15B-V1.0

 Text Generation • Updated 3 days ago • ↓ 12.5k • ❤ 332

CompVis/stable-diffusion-v1-4

 Text-to-Image • Updated about 17 hours ago • ↓ 448k • ❤ 5.72k

stabilityai/stable-diffusion-2-1

 Text-to-Image • Updated about 17 hours ago • ↓ 782k • ❤ 2.81k



Hugging Face

The AI community building the future.

[Verified](#)[Sponsor](#)[52.4k followers](#)[NYC + Paris](#)[https://huggingface.co/](#)[@huggingface](#)[Overview](#)[Repositories 339](#)[Projects 6](#)[Packages](#)[People 81](#)[Sponsoring 1](#)

Pinned

transformers [Public](#)

🤗 Transformers: the model-definition framework for state-of-the-art machine learning models in text, vision, audio, and multimodal models, for both inference and training.

[Python](#) [148k](#) [29.9k](#)

diffusers [Public](#)

🤗 Diffusers: State-of-the-art diffusion models for image, video, and audio generation in PyTorch and FLAX.

[Python](#) [30.2k](#) [6.2k](#)

datasets [Public](#)

🤗 The largest hub of ready-to-use datasets for AI models with fast, easy-to-use and efficient data manipulation tools

[Python](#) [20.5k](#) [2.9k](#)

peft [Public](#)

🤗 PEFT: State-of-the-art Parameter-Efficient Fine-Tuning.

[Python](#) [19.3k](#) [2k](#)

accelerate [Public](#)

optimum [Public](#)

People

[View all](#)

Sponsoring




Top languages

BERT: An influential NLP model

BERT: Bidirectional Encoder Representations from Transformers

Trained on:

- Wikipedia
- BooksCorps



BERT: An influential NLP model

“A visually stunning rumination on love”



A visually stunning rum##ation on love

BERT: An influential NLP model

“A visually stunning rumination on love”



<CLS> A visually stunning rum##ination on love <SEP>

BERT: An influential NLP model

“A visually stunning rumination on love”



<CLS> A visually stunning rumination on love <SEP>



101 1037 17453 14726 19379 12758 2006 2293 102

```
from transformers import BertTokenizerFast

sentences = [
    'I love my dog',
    'I love my cat'
]

local_tokenizer_path = "./bert_tokenizer_local"

# If loading directly from Hugging Face server, pass 'bert-base-uncased' as an argument
tokenizer = BertTokenizerFast.from_pretrained(local_tokenizer_path)

encoded_inputs = tokenizer(sentences, padding=True,
                           truncation=True, return_tensors='pt')

tokens = [tokenizer.convert_ids_to_tokens(ids)
          for ids in encoded_inputs["input_ids"]]

word_index = tokenizer.get_vocab() # For BertTokenizerFast, get_vocab() returns the vocab
```

```
from transformers import BertTokenizerFast

sentences = [
    'I love my dog',
    'I love my cat'
]

local_tokenizer_path = "./bert_tokenizer_local"

# If loading directly from Hugging Face server, pass 'bert-base-uncased' as an argument
tokenizer = BertTokenizerFast.from_pretrained(local_tokenizer_path)

encoded_inputs = tokenizer(sentences, padding=True,
                           truncation=True, return_tensors='pt')

tokens = [tokenizer.convert_ids_to_tokens(ids)
          for ids in encoded_inputs["input_ids"]]

word_index = tokenizer.get_vocab() # For BertTokenizerFast, get_vocab() returns the vocab
```

```
from transformers import BertTokenizerFast

sentences = [
    'I love my dog',
    'I love my cat'
]

local_tokenizer_path = "./bert_tokenizer_local"

# If loading directly from Hugging Face server, pass 'bert-base-uncased' as an argument
tokenizer = BertTokenizerFast.from_pretrained(local_tokenizer_path)

encoded_inputs = tokenizer(sentences, padding=True,
                           truncation=True, return_tensors='pt')

tokens = [tokenizer.convert_ids_to_tokens(ids)
          for ids in encoded_inputs["input_ids"]]

word_index = tokenizer.get_vocab() # For BertTokenizerFast, get_vocab() returns the vocab
```

```
from transformers import BertTokenizerFast

sentences = [
    'I love my dog',
    'I love my cat'
]

local_tokenizer_path = "./bert_tokenizer_local"

# If loading directly from Hugging Face server, pass 'bert-base-uncased' as an argument
tokenizer = BertTokenizerFast.from_pretrained(local_tokenizer_path)

encoded_inputs = tokenizer(sentences, padding=True,
                           truncation=True, return_tensors='pt')

tokens = [tokenizer.convert_ids_to_tokens(ids)
          for ids in encoded_inputs["input_ids"]]

word_index = tokenizer.get_vocab() # For BertTokenizerFast, get_vocab() returns the vocab
```

```
from transformers import BertTokenizerFast

sentences = [
    'I love my dog',
    'I love my cat'
]

local_tokenizer_path = "./bert_tokenizer_local"

# If loading directly from Hugging Face server, pass 'bert-base-uncased' as an argument
tokenizer = BertTokenizerFast.from_pretrained(local_tokenizer_path)

encoded_inputs = tokenizer(sentences, padding=True,
                           truncation=True, return_tensors='pt')

tokens = [tokenizer.convert_ids_to_tokens(ids)
          for ids in encoded_inputs["input_ids"]]

word_index = tokenizer.get_vocab() # For BertTokenizerFast, get_vocab() returns the vocab
```

```
from transformers import BertTokenizerFast

sentences = [
    'I love my dog',
    'I love my cat'
]

local_tokenizer_path = "./bert_tokenizer_local"

# If loading directly from Hugging Face server, pass 'bert-base-uncased' as an argument
tokenizer = BertTokenizerFast.from_pretrained(local_tokenizer_path)

encoded_inputs = tokenizer(sentences, padding=True,
                           truncation=True, return_tensors='pt')

tokens = [tokenizer.convert_ids_to_tokens(ids)
          for ids in encoded_inputs["input_ids"]]

word_index = tokenizer.get_vocab() # For BertTokenizerFast, get_vocab() returns the vocab
```

Using a pretrained tokenizer

```
# Print the human-readable `tokens` for each sentence
print("Tokens:", tokens)

print("\nToken IDs:", encoded_inputs['input_ids'])

# Print unique tokens from your sentences mapped to their unique IDs
helper_utils.print_unique_token_id_mappings(tokens, encoded_inputs['input_ids'])
```

Using a pretrained tokenizer

Output

```
Tokens: [[['[CLS]', 'i', 'love', 'my', 'dog', '[SEP]'], ['[CLS]', 'i', 'love', 'my', 'cat', '[SEP']']]  
Token IDs: tensor([[ 101, 1045, 2293, 2026, 3899, 102],  
                  [ 101, 1045, 2293, 2026, 4937, 102]])  
--- Unique Token to ID Mappings (for these sentences) ---  
[CLS] --> 101  
[SEP] --> 102  
i --> 1045  
my --> 2026  
love --> 2293  
dog --> 3899  
cat --> 4937
```



DeepLearning.AI

Tensorization

Working with text using PyTorch

Tensors represent batches of text




1D

Tensors represent batches of text



1D




2D


Tensors represent batches of text



1D



2D



3D

Tensorization

“I love my dog”



[1, 2, 3, 4]

```
import torch  
  
tokens = torch.tensor([1, 2, 3, 4])
```

Tensorization

“I love my dog”



[1, 2, 3, 4]

“I love my cat”



[1, 2, 3, 5]

“I like having pets”



[1, 6, 7, 8]

Tensorization

“I love my dog”



[1, 2, 3, 4]

“I love my cat”



[1, 2, 3, 5]

“I like having pets”




[1, 6, 7, 8]

[
[1, 2, 3, 4],
[1, 2, 3, 5],
[1, 6, 7, 8]
]

One-hot encoding

Words are represented as a vector with all zeros except in one position

One-hot encoding



Words are represented as a vector with all zeros except in one position

Bag of words: Count

Vocabulary:

ID:	'i '	'love '	'my '	'dog '	'cat '
	1	2	3	4	5



"I **love love** my dog"

[1 **2** 1 1 0]

"I love my cat"

[1 1 1 0 1]

Bag of words: Count

Vocabulary:

ID:

'i'	'love'	'my'	'dog'	'cat'
1	2	3	4	5



"I love love my dog"

[1	2	1	1	0]
----	---	---	---	----

"I love my cat"

[1	1	1	0	1]
------------	----------	----------	---	------------

"my cat I love"

[1	1	1	0	1]
------------	----------	----------	---	------------

Term frequency-inverse document frequency

TF



How often words appear in one document

IDF



Down-weights words that appear in many documents

$$TF(t, d) = \frac{\text{number of times } t \text{ appears in } d}{\text{total number of terms in } d}$$

$$IDF(t) = \log \frac{N}{1 + df}$$

$$TF - IDF(t, d) = TF(t, d) * IDF(t)$$

t: term

d: document

N: total number of documents

df: number of documents containing t

Term frequency-inverse document frequency

Vocabulary:

ID:

'i'	'love'	'my'	'dog'	'cat'
1	2	3	4	5



"I love love my dog"

[0	0	0	0.139	0]
----	---	---	-------	----

"I love my cat"

[0	0	0	0	0.173]
----	---	---	---	--------



DeepLearning.AI

Introduction to Embeddings

Working with text using PyTorch

Token ID's don't carry any meaning

“cat”

“dog”




[4937]




[3899]


Token ID's don't carry any meaning




Token ID's don't carry any meaning




Embeddings show how words are related




Embeddings show how words are related




Embeddings show how words are related




Embeddings show how words are related




Embeddings vs. one-hot vectors




Embeddings vs. one-hot vectors




Embeddings provide a map of meaning




Embeddings provide a map of meaning




Embeddings provide a map of meaning




Embeddings provide a map of meaning



Embeddings provide a map of meaning



Embeddings provide a map of meaning



Operations on embedding vectors

king - man

Operations on embedding vectors

king - man + woman

Operations on embedding vectors

$$\text{king} - \text{man} + \text{woman} \approx \text{queen}$$

Operations on embedding vectors

king = [4.2, 1.8, 7.5]

man = [2.1, 0.9, 5.0]

queen = [4.2, 2.1, 7.4]

woman = [2.0, 1.0, 5.1]

Operations on embedding vectors

king - man = [4.2 - 2.1, 1.8 - 0.9, 7.5 - 5.0]

= [2.1, 0.9, 2.5]

Operations on embedding vectors

$$\begin{aligned} \text{king} - \text{man} + \text{woman} &= [2.1 + 2.0, 0.9 + 1.0, 2.5 + 5.1] \\ &= [4.1, 1.9, 7.6] \end{aligned}$$

Operations on embedding vectors


king - man + woman = [2.1 + 2.0, 0.9 + 1.0, 2.5 + 5.1]

$$= [4.1, 1.9, 7.6]$$




queen = [4.2, 2.1, 7.4]

Distributional hypothesis




Distributional hypothesis




Distributional hypothesis




Distributional hypothesis




Cosine similarity



Cosine ≈ 1
Similar




Cosine ≈ 0
Orthogonal




Cosine ≈ -1
Opposite

$$\text{similarity}(A, B) = \cos(\theta) = \frac{A \cdot B}{\|A\| \|B\|}$$


Cosine similarity




Cosine similarity




Embedding layers




Embedding layers



Embedding layers



Embedding layers



Static embeddings

They map each word in the vocabulary to one fixed vector

Static embeddings

They map each word in the vocabulary to one fixed vector



“flying **bat**”



“baseball **bat**”



bat = [0.32, 2.24, 6.30, 0.76]

Static embedding models

Word2Vec

GloVe

FastText

Contextual embeddings

They assign embedding vectors based on context

Contextual embeddings

They assign embedding vectors based on context



“flying **bat**”



bat = [0.12, 2.65, 6.27, 0.56]



“baseball **bat**”



bat = [0.09, 2.51, 6.90, 0.42]

Contextual embedding models

BERT

GPT

ELMo

Trade-offs with contextual embeddings



Pros

- Catch subtle differences
- Good for sentiment analysis, NER and translation

Trade-offs with contextual embeddings



Pros

- Catch subtle differences
- Good for sentiment analysis, NER and translation



Cons

- Computationally heavier
- Compute new embeddings for every sentence



DeepLearning.AI

Implementing Embeddings in PyTorch

Working with text using PyTorch

GloVe: A static embedding model



Stands for Global Vectors for Word Representation

GloVe: A static embedding model



Stands for Global Vectors for Word Representation



Learns by looking at how often words appear together

GloVe: A static embedding model



Stands for Global Vectors for Word Representation



Learns by looking at how often words appear together



GloVe 6B: 6 billion tokens from Wikipedia and Gigaword

Using GloVe in PyTorch

```
# Download the data for the GloVe 6B 100d model
helper_utils.download_glove6B()

# Specify the path to the 100d GloVe file
glove_file = './glove_data/glove.6B.100d.txt'

# Load the pre-trained word vectors from the file
glove_embeddings = helper_utils.load_glove_embeddings(glove_file)
```

Using GloVe in PyTorch

```
# Download the data for the GloVe 6B 100d model
helper_utils.download_glove6B()
```

```
# Specify the path to the 100d GloVe file
glove_file = './glove_data/glove.6B.100d.txt'
```

```
# Load the pre-trained word vectors from the file
glove_embeddings = helper_utils.load_glove_embeddings(glove_file)
```

Using GloVe in PyTorch

```
# Download the data for the GloVe 6B 100d model
helper_utils.download_glove6B()

# Specify the path to the 100d GloVe file
glove_file = './glove_data/glove.6B.100d.txt'

# Load the pre-trained word vectors from the file
glove_embeddings = helper_utils.load_glove_embeddings(glove_file)
```

```
def find_closest_words(embedding, embeddings_dict, exclude_words=[], top_n=5):

    filtered_words = [word for word in embeddings_dict.keys() if word not in exclude_words]

    if not filtered_words:
        return None

    embedding_matrix = np.array([embeddings_dict[word] for word in filtered_words])

    target_embedding = embedding.reshape(1, -1)

    similarity_scores = cosine_similarity(target_embedding, embedding_matrix)

    closest_word_indices = np.argsort(similarity_scores[0])[:-1][:top_n]

    return [(filtered_words[i], similarity_scores[0][i]) for i in closest_word_indices]
```

```
def find_closest_words(embedding, embeddings_dict, exclude_words=[], top_n=5):  
    filtered_words = [word for word in embeddings_dict.keys() if word not in exclude_words]  
  
    if not filtered_words:  
        return None  
  
    embedding_matrix = np.array([embeddings_dict[word] for word in filtered_words])  
  
    target_embedding = embedding.reshape(1, -1)  
  
    similarity_scores = cosine_similarity(target_embedding, embedding_matrix)  
  
    closest_word_indices = np.argsort(similarity_scores[0])[::-1][:top_n]  
  
    return [(filtered_words[i], similarity_scores[0][i]) for i in closest_word_indices]
```

```
def find_closest_words(embedding, embeddings_dict, exclude_words=[], top_n=5):

    filtered_words = [word for word in embeddings_dict.keys() if word not in exclude_words]

    if not filtered_words:
        return None

    embedding_matrix = np.array([embeddings_dict[word] for word in filtered_words])

    target_embedding = embedding.reshape(1, -1)

    similarity_scores = cosine_similarity(target_embedding, embedding_matrix)

    closest_word_indices = np.argsort(similarity_scores[0])[:-1][:top_n]

    return [(filtered_words[i], similarity_scores[0][i]) for i in closest_word_indices]
```

```
def find_closest_words(embedding, embeddings_dict, exclude_words=[], top_n=5):  
    filtered_words = [word for word in embeddings_dict.keys() if word not in exclude_words]  
  
    if not filtered_words:  
        return None  
  
    embedding_matrix = np.array([embeddings_dict[word] for word in filtered_words])  
  
    target_embedding = embedding.reshape(1, -1)  
  
    similarity_scores = cosine_similarity(target_embedding, embedding_matrix)  
  
    closest_word_indices = np.argsort(similarity_scores[0])[::-1][:top_n]  
  
    return [(filtered_words[i], similarity_scores[0][i]) for i in closest_word_indices]
```

Using GloVe in PyTorch

```
if all(word in glove_embeddings for word in ['king', 'man', 'woman']):  
    king = glove_embeddings['king']  
    man = glove_embeddings['man']  
    woman = glove_embeddings['woman']  
  
result_embedding = king - man + woman  
  
top_n = 5  
  
closest_words_with_scores = find_closest_words(  
    result_embedding,  
    glove_embeddings,  
    exclude_words=['king', 'man', 'woman'],  
    top_n=top_n  
)
```

Using GloVe in PyTorch

```
if all(word in glove_embeddings for word in ['king', 'man', 'woman']):  
    king = glove_embeddings['king']  
    man = glove_embeddings['man']  
    woman = glove_embeddings['woman']  
  
result_embedding = king - man + woman  
  
top_n = 5  
  
closest_words_with_scores = find_closest_words(  
    result_embedding,  
    glove_embeddings,  
    exclude_words=['king', 'man', 'woman'],  
    top_n=top_n  
)
```

Using GloVe in PyTorch

```
if all(word in glove_embeddings for word in ['king', 'man', 'woman']):  
    king = glove_embeddings['king']  
    man = glove_embeddings['man']  
    woman = glove_embeddings['woman']  
  
    result_embedding = king - man + woman  
  
    top_n = 5  
  
    closest_words_with_scores = find_closest_words(  
        result_embedding,  
        glove_embeddings,  
        exclude_words=['king', 'man', 'woman'],  
        top_n=top_n  
    )
```

Using GloVe in PyTorch

```
if all(word in glove_embeddings for word in ['king', 'man', 'woman']):  
    king = glove_embeddings['king']  
    man = glove_embeddings['man']  
    woman = glove_embeddings['woman']  
  
    result_embedding = king - man + woman  
  
    top_n = 5  
  
    closest_words_with_scores = find_closest_words(  
        result_embedding,  
        glove_embeddings,  
        exclude_words=['king', 'man', 'woman'],  
        top_n=top_n  
    )
```

Using GloVe in PyTorch

Output

```
king - man + woman ≈ queen (Score: 0.7834)
```

```
--- Other Top 4 Results ---
```

```
monarch (Score: 0.6934)
```

```
throne (Score: 0.6833)
```

```
daughter (Score: 0.6809)
```

```
prince (Score: 0.6713)
```

Static embeddings can't handle polysemy

"A **bat** flew out of the cave"



Static embeddings can't handle polysemy

"A **bat** flew out of the cave"



"He swung the baseball **bat**"



GloVe produces the same vector for 'bat'

```
sentence1 = "A bat flew out of the cave."  
sentence2 = "He swung the baseball bat."  
  
bat_from_sentence1 = glove_embeddings["bat"]  
bat_from_sentence2 = glove_embeddings["bat"]  
  
are_identical = np.array_equal(bat_from_sentence1, bat_from_sentence2)  
print(f"Are the vectors for 'bat' from each sentence identical? {are_identical}")
```

GloVe produces the same vector for 'bat'

```
sentence1 = "A bat flew out of the cave."  
sentence2 = "He swung the baseball bat."  
  
bat_from_sentence1 = glove_embeddings["bat"]  
bat_from_sentence2 = glove_embeddings["bat"]  
  
are_identical = np.array_equal(bat_from_sentence1, bat_from_sentence2)  
print(f"Are the vectors for 'bat' from each sentence identical? {are_identical}")
```

GloVe produces the same vector for 'bat'

```
sentence1 = "A bat flew out of the cave."  
sentence2 = "He swung the baseball bat."  
  
bat_from_sentence1 = glove_embeddings["bat"]  
bat_from_sentence2 = glove_embeddings["bat"]  
  
are_identical = np.array_equal(bat_from_sentence1, bat_from_sentence2)  
print(f"Are the vectors for 'bat' from each sentence identical? {are_identical}")
```

GloVe produces the same vector for 'bat'

```
sentence1 = "A bat flew out of the cave."  
sentence2 = "He swung the baseball bat."  
  
bat_from_sentence1 = glove_embeddings["bat"]  
bat_from_sentence2 = glove_embeddings["bat"]  
  
are_identical = np.array_equal(bat_from_sentence1, bat_from_sentence2)  
print(f"Are the vectors for 'bat' from each sentence identical? {are_identical}")
```

Output

Are the vectors for 'bat' from each sentence identical? True

Using a contextual embedding model: BERT

```
helper_utils.download_bert()
bert_path = './bert_model'
tokenizer, model_bert = helper_utils.load_bert(bert_path)
```

Using a contextual embedding model: BERT

```
# --- Process and Print Vectors for Sentence 1 ---
print("--- Sentence 1 (first 5 values) ---")
inputs1 = tokenizer(sentence1, return_tensors='pt')
with torch.no_grad():
    outputs1 = model_bert(**inputs1)
last_hidden_state1 = outputs1.last_hidden_state[0] # Embeddings for all tokens

tokens1 = tokenizer.convert_ids_to_tokens(inputs1['input_ids'][0])
```

Using a contextual embedding model: BERT

```
# --- Process and Print Vectors for Sentence 1 ---
print("--- Sentence 1 (first 5 values) ---")
inputs1 = tokenizer(sentence1, return_tensors='pt')
with torch.no_grad():
    outputs1 = model_bert(**inputs1)
last_hidden_state1 = outputs1.last_hidden_state[0] # Embeddings for all tokens

tokens1 = tokenizer.convert_ids_to_tokens(inputs1['input_ids'][0])
```

Using a contextual embedding model: BERT

```
# --- Process and Print Vectors for Sentence 1 ---
print("--- Sentence 1 (first 5 values) ---")
inputs1 = tokenizer(sentence1, return_tensors='pt')

with torch.no_grad():
    outputs1 = model_bert(**inputs1)
last_hidden_state1 = outputs1.last_hidden_state[0] # Embeddings for all tokens

tokens1 = tokenizer.convert_ids_to_tokens(inputs1['input_ids'][0])
```

Using a contextual embedding model: BERT

```
# --- Process and Print Vectors for Sentence 1 ---
print("--- Sentence 1 (first 5 values) ---")
inputs1 = tokenizer(sentence1, return_tensors='pt')
with torch.no_grad():
    outputs1 = model_bert(**inputs1)
last_hidden_state1 = outputs1.last_hidden_state[0] # Embeddings for all tokens

tokens1 = tokenizer.convert_ids_to_tokens(inputs1['input_ids'][0])
```

Using a contextual embedding model: BERT

```
# --- Process and Print Vectors for Sentence 2 ---
print("--- Sentence 2 (first 5 values) ---")
inputs2 = tokenizer(sentence2, return_tensors='pt')
with torch.no_grad():
    outputs2 = model_bert(**inputs2)
last_hidden_state2 = outputs2.last_hidden_state[0] # Embeddings for all tokens

tokens2 = tokenizer.convert_ids_to_tokens(inputs2['input_ids'][0])
```

BERT produces contextual embeddings

```
bat_animal_vector = last_hidden_state1[2].numpy()
bat_sport_vector = last_hidden_state2[5].numpy()
are_identical = np.array_equal(bat_animal_vector, bat_sport_vector)
print(f"Are the contextual BERT vectors for 'bat' identical? {are_identical}")
```

BERT produces contextual embeddings

```
bat_animal_vector = last_hidden_state1[2].numpy()
bat_sport_vector = last_hidden_state2[5].numpy()
are_identical = np.array_equal(bat_animal_vector, bat_sport_vector)
print(f"Are the contextual BERT vectors for 'bat' identical? {are_identical}")
```

Output

```
Are the contextual BERT vectors for 'bat' identical? False
```

When to use static vs. contextual?

Static embeddings



Simple tasks where
speed and memory
efficiency are prioritized

When to use static vs. contextual?

Static embeddings



Simple tasks where speed and memory efficiency are prioritized

Contextual embeddings



Advanced tasks where subtle differences are critical

Visualization techniques for embeddings

PCA
Principal Component
Analysis

Efficient, but can distort local
relationships

Visualization techniques for embeddings

PCA

Principal Component Analysis

Efficient, but can distort local relationships

t-SNE


t-distributed Stochastic Neighbor Embedding

More intensive, but preserves local structure and clustering

Visualizing word embeddings




Visualizing word embeddings




Koehrsen, 2018

Visualizing word embeddings



Koehrsen, 2018

Visualizing word embeddings



Koehrsen, 2018

Embeddings visualization in PyTorch

```
words_to_visualize = ['car', 'bike', 'plane',      # Category: Vehicles
                      'cat', 'dog', 'bird',       # Category: Pets
                      'orange', 'apple', 'grape' # Category: Fruits
]

visualization_dict = {
    'Vehicle': ['car', 'bike', 'plane'],
    'Pet': ['cat', 'dog', 'bird'],
    'Fruit': ['orange', 'apple', 'grape']
}
```

Embeddings visualization in PyTorch

```
embedding_vectors_list = []

for word in words_to_visualize:
    embedding_vectors_list.append(glove_embeddings[word])

embedding_vectors = np.array(embedding_vectors_list)

reducer = PCA(n_components=2)

coords_2d = reducer.fit_transform(embedding_vectors)

helper_utils.plot_embeddings(coords=coords_2d,
                             labels=words_to_visualize,
                             label_dict=visualization_dict,
                             title='GloVe Pre-Trained Embeddings'
                             )
```

Embeddings visualization in PyTorch

```
embedding_vectors_list = []

for word in words_to_visualize:
    embedding_vectors_list.append(glove_embeddings[word])

embedding_vectors = np.array(embedding_vectors_list)

reducer = PCA(n_components=2)

coords_2d = reducer.fit_transform(embedding_vectors)

helper_utils.plot_embeddings(coords=coords_2d,
                             labels=words_to_visualize,
                             label_dict=visualization_dict,
                             title='GloVe Pre-Trained Embeddings'
                             )
```

Embeddings visualization in PyTorch

```
embedding_vectors_list = []

for word in words_to_visualize:
    embedding_vectors_list.append(glove_embeddings[word])


embedding_vectors = np.array(embedding_vectors_list)

reducer = PCA(n_components=2)

coords_2d = reducer.fit_transform(embedding_vectors)

helper_utils.plot_embeddings(coords=coords_2d,
                             labels=words_to_visualize,
                             label_dict=visualization_dict,
                             title='GloVe Pre-Trained Embeddings'
                             )
```

Embeddings visualization in PyTorch



Practical considerations about embeddings

Practical considerations about embeddings



Dimensionality

Practical considerations about embeddings



Dimensionality



Initialization

Practical considerations about embeddings



Dimensionality



Initialization



Out of
vocabulary
problem

Practical considerations about embeddings



Dimensionality



Initialization



Out of
vocabulary
problem



Type of task



DeepLearning.AI

Building a Simple Text Classifier in PyTorch

Working with text using PyTorch



Gmail

Compose

Inbox 152

Starred

Snoozed

Sent

Drafts

Search mail

Primary


Striving Blogger

Carly - Blogging Li.


Tyler at ConvertKit

1



This image shows the Gmail mobile application interface. At the top, there's a search bar labeled "Search mail". Below it are standard navigation icons: a square with a downward arrow, a reply arrow, a refresh symbol, and a more options menu. The main content area is divided into sections by red horizontal lines. The first section is "Primary", which contains three email previews. The second section is "Inbox", which has a red notification badge showing "152". Other sections listed are "Starred", "Snoozed", "Sent", and "Drafts". On the far left, there's a "Compose" button with a plus sign icon. The bottom of the screen features a footer with the number "1".




Text classification pipeline



Text classification pipeline




Text classification pipeline



Pooling techniques


Mean pooling



Takes the average
across all
embeddings

Pooling techniques

Mean pooling



Takes the average
across all
embeddings


Max pooling



Finds the dominant
feature across each
dimension

Pooling techniques

Mean pooling




Takes the average
across all
embeddings

Max pooling




Finds the dominant
feature across each
dimension

Sum pooling




Adds embeddings
together


Text classification pipeline




Text classification pipeline



nn.EmbeddingBag




nn.EmbeddingBag



Why use it?

- Better performance
- Simplicity

`nn.EmbeddingBag`



Why use it?

- Better performance
- Simplicity

How to use it?

- Provide token IDs and an offset tensor
- Pools all embeddings in the chosen mode

Building a text classifier: Preprocessing



Select and load a dataset



Clean text



Tokenize



Split training vs. validation set



Build a vocabulary



Create DataLoader

The dataset: Food.com recipe collection

Fruit-based



“Apple a Day Milkshake”

Vegetable-based



“Zuppa Toscana”

Build a classifier based on different models

Type of embedding

EmbeddingBag

Manual pooling

Type of pooling

Mean

Mean

Max

Sum

Building a text classifier with EmbeddingBag

```
class EmbeddingBagClassifier(nn.Module):

    def __init__(self, vocab_size, embedding_dim, num_classes):
        super().__init__()
        self.embedding_bag = nn.EmbeddingBag(vocab_size, embedding_dim, mode='mean')
        self.dropout = nn.Dropout(0.5)
        self.fc = nn.Linear(embedding_dim, num_classes)

    def forward(self, text, offsets=None):
        embedded = self.embedding_bag(text, offsets)
        embedded = self.dropout(embedded)
        return self.fc(embedded)
```

Building a text classifier with EmbeddingBag

```
class EmbeddingBagClassifier(nn.Module):

    def __init__(self, vocab_size, embedding_dim, num_classes):
        super().__init__()
        self.embedding_bag = nn.EmbeddingBag(vocab_size, embedding_dim, mode='mean')
        self.dropout = nn.Dropout(0.5)
        self.fc = nn.Linear(embedding_dim, num_classes)

    def forward(self, text, offsets=None):
        embedded = self.embedding_bag(text, offsets)
        embedded = self.dropout(embedded)
        return self.fc(embedded)
```

Building a text classifier with EmbeddingBag

```
class EmbeddingBagClassifier(nn.Module):

    def __init__(self, vocab_size, embedding_dim, num_classes):
        super().__init__()
        self.embedding_bag = nn.EmbeddingBag(vocab_size, embedding_dim, mode='mean')
        self.dropout = nn.Dropout(0.5)
        self.fc = nn.Linear(embedding_dim, num_classes)

    def forward(self, text, offsets=None):
        embedded = self.embedding_bag(text, offsets)
        embedded = self.dropout(embedded)
        return self.fc(embedded)
```

Building a text classifier with EmbeddingBag

```
class EmbeddingBagClassifier(nn.Module):

    def __init__(self, vocab_size, embedding_dim, num_classes):
        super().__init__()
        self.embedding_bag = nn.EmbeddingBag(vocab_size, embedding_dim, mode='mean')
        self.dropout = nn.Dropout(0.5)
        self.fc = nn.Linear(embedding_dim, num_classes)

    def forward(self, text, offsets=None):
        embedded = self.embedding_bag(text, offsets)
        embedded = self.dropout(embedded)
        return self.fc(embedded)
```

Building a text classifier with EmbeddingBag

```
class EmbeddingBagClassifier(nn.Module):

    def __init__(self, vocab_size, embedding_dim, num_classes):
        super().__init__()
        self.embedding_bag = nn.EmbeddingBag(vocab_size, embedding_dim, mode='mean')
        self.dropout = nn.Dropout(0.5)
        self.fc = nn.Linear(embedding_dim, num_classes)

    def forward(self, text, offsets=None):
        embedded = self.embedding_bag(text, offsets)
        embedded = self.dropout(embedded)
        return self.fc(embedded)
```

Building a text classifier with manual pooling

```
class ManualPoolingClassifier(nn.Module):

    def __init__(self, vocab_size, embedding_dim, num_classes, pooling='mean'):
        super().__init__()
        self.embedding = nn.Embedding(vocab_size, embedding_dim, padding_idx=0)
        self.pooling = pooling
        self.fc = nn.Linear(embedding_dim, num_classes)
        self.dropout = nn.Dropout(0.5)
```

Building a text classifier with manual pooling

```
class ManualPoolingClassifier(nn.Module):  
  
    def __init__(self, vocab_size, embedding_dim, num_classes, pooling='mean'):  
        super().__init__()  
        self.embedding = nn.Embedding(vocab_size, embedding_dim, padding_idx=0)  
        self.pooling = pooling  
        self.fc = nn.Linear(embedding_dim, num_classes)  
        self.dropout = nn.Dropout(0.5)
```

Building a text classifier with manual pooling

```
class ManualPoolingClassifier(nn.Module):  
  
    def __init__(self, vocab_size, embedding_dim, num_classes, pooling='mean'):  
        super().__init__()  
        self.embedding = nn.Embedding(vocab_size, embedding_dim, padding_idx=0)  
        self.pooling = pooling  
        self.fc = nn.Linear(embedding_dim, num_classes)  
        self.dropout = nn.Dropout(0.5)
```

Building a text classifier with manual pooling

```
class ManualPoolingClassifier(nn.Module):

    def __init__(self, vocab_size, embedding_dim, num_classes, pooling='mean'):
        super().__init__()
        self.embedding = nn.Embedding(vocab_size, embedding_dim, padding_idx=0)
        self.pooling = pooling
        self.fc = nn.Linear(embedding_dim, num_classes)
        self.dropout = nn.Dropout(0.5)
```

Building a text classifier with manual pooling

```
def forward(self, text):
    embedded = self.embedding(text)
    mask = (text != 0).float().unsqueeze(-1)
    embedded = embedded * mask

    if self.pooling == 'mean':
        pooled = embedded.sum(dim=1) / mask.sum(dim=1).clamp(min=1)
    elif self.pooling == 'max':
        embedded[mask.squeeze(-1) == 0] = float('-inf')
        pooled, _ = embedded.max(dim=1)
    elif self.pooling == 'sum':
        pooled = embedded.sum(dim=1)

    pooled = self.dropout(pooled)
    return self.fc(pooled)
```

Building a text classifier with manual pooling

```
def forward(self, text):
    embedded = self.embedding(text)
    mask = (text != 0).float().unsqueeze(-1)
    embedded = embedded * mask

    if self.pooling == 'mean':
        pooled = embedded.sum(dim=1) / mask.sum(dim=1).clamp(min=1)
    elif self.pooling == 'max':
        embedded[mask.squeeze(-1) == 0] = float('-inf')
        pooled, _ = embedded.max(dim=1)
    elif self.pooling == 'sum':
        pooled = embedded.sum(dim=1)

    pooled = self.dropout(pooled)
    return self.fc(pooled)
```

Building a text classifier with manual pooling

```
def forward(self, text):
    embedded = self.embedding(text)
    mask = (text != 0).float().unsqueeze(-1)
    embedded = embedded * mask

    if self.pooling == 'mean':
        pooled = embedded.sum(dim=1) / mask.sum(dim=1).clamp(min=1)
    elif self.pooling == 'max':
        embedded[mask.squeeze(-1) == 0] = float('-inf')
        pooled, _ = embedded.max(dim=1)
    elif self.pooling == 'sum':
        pooled = embedded.sum(dim=1)

    pooled = self.dropout(pooled)
    return self.fc(pooled)
```

Building a text classifier with manual pooling

```
def forward(self, text):
    embedded = self.embedding(text)
    mask = (text != 0).float().unsqueeze(-1)
    embedded = embedded * mask

    if self.pooling == 'mean':
        pooled = embedded.sum(dim=1) / mask.sum(dim=1).clamp(min=1)
    elif self.pooling == 'max':
        embedded[mask.squeeze(-1) == 0] = float('-inf')
        pooled, _ = embedded.max(dim=1)
    elif self.pooling == 'sum':
        pooled = embedded.sum(dim=1)

    pooled = self.dropout(pooled)
    return self.fc(pooled)
```

Building a text classifier with manual pooling

```
def forward(self, text):
    embedded = self.embedding(text)
    mask = (text != 0).float().unsqueeze(-1)
    embedded = embedded * mask

    if self.pooling == 'mean':
        pooled = embedded.sum(dim=1) / mask.sum(dim=1).clamp(min=1)
    elif self.pooling == 'max':
        embedded[mask.squeeze(-1) == 0] = float('-inf')
        pooled, _ = embedded.max(dim=1)
    elif self.pooling == 'sum':
        pooled = embedded.sum(dim=1)

    pooled = self.dropout(pooled)
    return self.fc(pooled)
```

Building a text classifier with manual pooling

```
def forward(self, text):
    embedded = self.embedding(text)
    mask = (text != 0).float().unsqueeze(-1)
    embedded = embedded * mask

    if self.pooling == 'mean':
        pooled = embedded.sum(dim=1) / mask.sum(dim=1).clamp(min=1)
    elif self.pooling == 'max':
        embedded[mask.squeeze(-1) == 0] = float('-inf')
        pooled, _ = embedded.max(dim=1)
    elif self.pooling == 'sum':
        pooled = embedded.sum(dim=1)

    pooled = self.dropout(pooled)
    return self.fc(pooled)
```

Building a text classifier with manual pooling

```
vocab_size = len(vocab)
embedding_dim = 64
num_classes = 2

model_embag = EmbeddingBagClassifier(vocab_size, embedding_dim, num_classes)
model_manual_mean = ManualPoolingClassifier(vocab_size, embedding_dim,
num_classes, pooling='mean')
model_manual_max = ManualPoolingClassifier(vocab_size, embedding_dim,
num_classes, pooling='max')
model_manual_sum = ManualPoolingClassifier(vocab_size, embedding_dim,
num_classes, pooling='sum')
```

Building a text classifier with manual pooling

```
vocab_size = len(vocab)
embedding_dim = 64
num_classes = 2
```

```
model_embag = EmbeddingBagClassifier(vocab_size, embedding_dim, num_classes)
model_manual_mean = ManualPoolingClassifier(vocab_size, embedding_dim,
num_classes, pooling='mean')
model_manual_max = ManualPoolingClassifier(vocab_size, embedding_dim,
num_classes, pooling='max')
model_manual_sum = ManualPoolingClassifier(vocab_size, embedding_dim,
num_classes, pooling='sum')
```

Building a text classifier with manual pooling

```
vocab_size = len(vocab)
embedding_dim = 64
num_classes = 2

model_embag = EmbeddingBagClassifier(vocab_size, embedding_dim, num_classes)
model_manual_mean = ManualPoolingClassifier(vocab_size, embedding_dim,
num_classes, pooling='mean')
model_manual_max = ManualPoolingClassifier(vocab_size, embedding_dim,
num_classes, pooling='max')
model_manual_sum = ManualPoolingClassifier(vocab_size, embedding_dim,
num_classes, pooling='sum')
```



DeepLearning.AI

Fine Tuning Pretrained Text Classification Models

Working with text using PyTorch

Why use pretrained models?



Preloaded
with semantic
knowledge

Why use pretrained models?



Preloaded
with semantic
knowledge



Transfer
learning
benefits

Why use pretrained models?



Preloaded
with semantic
knowledge



Transfer
learning
benefits



Efficiency

Why use pretrained models?




Preloaded
with semantic
knowledge



Transfer
learning
benefits




Efficiency




Better
generalization

DistilBERT


- Transformer model
- Use cases:
 - Text classification
 - Named entity recognition
 - Question answering, and more



Transfer learning



Fine tuning



Considerations when fine-tuning



Advantages

- Adapt to specific vocab
- Optimize for specific goal
- Higher performance

Considerations when fine-tuning



Advantages

- Adapt to specific vocab
- Optimize for specific goal
- Higher performance



Risks

- Catastrophic forgetting
- Overfitting
- Training instability

Loading DISTILBert

```
model_name="distilbert-base-uncased"
model_path="./distilbert-local-base"

# Ensure the model is downloaded
helper_utils.download_bert(model_name, model_path)

bert_model, bert_tokenizer = helper_utils.load_bert(model_path, num_classes=2)
```

Loading DISTILBert

```
model_name="distilbert-base-uncased"
model_path="./distilbert-local-base"

# Ensure the model is downloaded
helper_utils.download_bert(model_name, model_path)

bert_model, bert_tokenizer = helper_utils.load_bert(model_path, num_classes=2)
```

Partial fine-tuning with DistilBERT

```
for param in bert_model.parameters():
    param.requires_grad = False

layers_to_train = 2
transformer_layers = bert_model.distilbert.transformer.layer
for i in range(layers_to_train):
    layer_to_unfreeze = transformer_layers[-(i+1)]

    for param in layer_to_unfreeze.parameters():
        param.requires_grad = True

for param in bert_model.pre_classifier.parameters():
    param.requires_grad = True

for param in bert_model.classifier.parameters():
    param.requires_grad = True
```

Partial fine-tuning with DistilBERT

```
for param in bert_model.parameters():
    param.requires_grad = False

layers_to_train = 2
transformer_layers = bert_model.distilbert.transformer.layer
for i in range(layers_to_train):
    layer_to_unfreeze = transformer_layers[-(i+1)]

    for param in layer_to_unfreeze.parameters():
        param.requires_grad = True

for param in bert_model.pre_classifier.parameters():
    param.requires_grad = True

for param in bert_model.classifier.parameters():
    param.requires_grad = True
```

Partial fine-tuning with DistilBERT

```
for param in bert_model.parameters():
    param.requires_grad = False

layers_to_train = 2
transformer_layers = bert_model.distilbert.transformer.layer
for i in range(layers_to_train):
    layer_to_unfreeze = transformer_layers[-(i+1)]

    for param in layer_to_unfreeze.parameters():
        param.requires_grad = True

for param in bert_model.pre_classifier.parameters():
    param.requires_grad = True

for param in bert_model.classifier.parameters():
    param.requires_grad = True
```

Partial fine-tuning with DistilBERT

```
for param in bert_model.parameters():
    param.requires_grad = False

layers_to_train = 2
transformer_layers = bert_model.distilbert.transformer.layer
for i in range(layers_to_train):
    layer_to_unfreeze = transformer_layers[-(i+1)]

    for param in layer_to_unfreeze.parameters():
        param.requires_grad = True

for param in bert_model.pre_classifier.parameters():
    param.requires_grad = True

for param in bert_model.classifier.parameters():
    param.requires_grad = True
```

Partial fine-tuning with DistilBERT

```
for param in bert_model.parameters():
    param.requires_grad = False

layers_to_train = 2
transformer_layers = bert_model.distilbert.transformer.layer
for i in range(layers_to_train):
    layer_to_unfreeze = transformer_layers[-(i+1)]

    for param in layer_to_unfreeze.parameters():
        param.requires_grad = True

for param in bert_model.pre_classifier.parameters():
    param.requires_grad = True

for param in bert_model.classifier.parameters():
    param.requires_grad = True
```

Partial fine-tuning with DistilBERT

```
for param in bert_model.parameters():
    param.requires_grad = False

layers_to_train = 2
transformer_layers = bert_model.distilbert.transformer.layer
for i in range(layers_to_train):
    layer_to_unfreeze = transformer_layers[-(i+1)]

    for param in layer_to_unfreeze.parameters():
        param.requires_grad = True

for param in bert_model.pre_classifier.parameters():
    param.requires_grad = True

for param in bert_model.classifier.parameters():
    param.requires_grad = True
```

Partial fine-tuning with DistilBERT

```
for param in bert_model.parameters():
    param.requires_grad = False

layers_to_train = 2
transformer_layers = bert_model.distilbert.transformer.layer
for i in range(layers_to_train):
    layer_to_unfreeze = transformer_layers[-(i+1)]
    for param in layer_to_unfreeze.parameters():
        param.requires_grad = True

for param in bert_model.pre_classifier.parameters():
    param.requires_grad = True

for param in bert_model.classifier.parameters():
    param.requires_grad = True
```



Every architecture is different!

What you learned in this module



Unique challenges with text data

What you learned in this module



Unique challenges with text data



Tokenization and tensorization

What you learned in this module



Unique challenges with text data



Tokenization and tensorization



Embeddings: turning words into vectors

What you learned in this module



Unique challenges with text data



Tokenization and tensorization



Embeddings: turning words into vectors



Building a text classification pipeline

What you learned in this module



Unique challenges with text data



Tokenization and tensorization



Embeddings: turning words into vectors



Building a text classification pipeline



Using pretrained models and fine-tuning